Analyze A/B Test Results

This project will assure you have mastered the subjects covered in the statistics lessons. The hope is to have this project be as comprehensive of these topics as possible. Good luck!

Table of Contents

Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these

For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

As you work through this notebook, follow along in the classroom and answer the corresponding quiz questions associated with each question. The labels for each classroom concept are provided for each question. This will assure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the criteria. As a final check, assure you meet all the criteria on the RUBRIC.

Part I - Probability

To get started, let's import our libraries.

In [1]:
```
import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline

import seaborn as sns
sns.set_style('darkgrid')

#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

`1.` Now, read in the `ab_data.csv` data. Store it in `df`. **Use your dataframe to answer the questions in Quiz 1 of the classroom.**

a. Read in the dataset and take a look at the top few rows here:

In [2]:
```
df = pd.read_csv('ab_data.csv')
df.head()
```

Out[2]:

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

b. Use the below cell to find the number of rows in the dataset.

In [3]:
```
df.shape[0]
```

Out[3]:
```
294478
```

c. The number of unique users in the dataset.

In [4]:
```
df['user_id'].nunique()
```

Out[4]:
```
290584
```

d. The proportion of users converted.

In [5]:
```
converted = df.query('converted == 1')
converted.shape[0] /df.shape[0]
```

Out[5]:
```
0.11965919355605512
```

e. The number of times the `new_page` and `treatment` don't line up.

In [6]:
```
df.query('group == "treatment" and landing_page != "new_page"').shape[0] +\
df.query('group != "treatment" and landing_page == "new_page"').shape[0]
```

Out[6]:
```
3893
```

f. Do any of the rows have missing values?

In [7]:
```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
 #   Column        Non-Null Count   Dtype 
---  ------        --------------   ----- 
 0   user_id       294478 non-null  int64 
 1   timestamp     294478 non-null  object
 2   group         294478 non-null  object
 3   landing_page  294478 non-null  object
 4   converted     294478 non-null  int64 
dtypes: int64(2), object(3)
memory usage: 11.2+ MB
```

No, there is'nt missing values.

`2.` For the rows where **treatment** is not aligned with **new_page** or **control** is not aligned with **old_page**, we cannot be sure if this row truly received the new or old page. Use **Quiz 2** in the classroom to provide how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in **df2**.

In [8]:
```
df2=df
df2.drop(df.query("(group == 'treatment' and landing_page == 'old_page') or (group ==
'control' and landing_page == 'new_page')").index, inplace=True)
```

In [9]:
```
# Double Check all of the correct rows were removed - this should be 0
df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) ==
False].shape[0]
```

Out[9]:
```
0
```

`3.` Use **df2** and the cells below to answer questions for **Quiz3** in the classroom.

a. How many unique **user_id**s are in **df2**?

In [10]:
```
df2['user_id'].nunique()
```

Out[10]:
```
290584
```

b. There is one **user_id** repeated in **df2**. What is it?

In [11]:
```
df2[df2.duplicated(['user_id'], keep=False)]['user_id'].unique()
```

Out[11]:
```
array([773192])
```

c. What is the row information for the repeat **user_id**?

In [12]:
```
df2[df2.duplicated(['user_id'], keep=False)]
```

Out[12]:

	user_id	timestamp	group	landing_page	converted
1899	773192	2017-01-09 05:37:58.781806	treatment	new_page	0
2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

d. Remove **one** of the rows with a duplicate **user_id**, but keep your dataframe as **df2**.

In [13]:
```
df2.drop(df2.loc[[1899]].index, inplace=True)
```

`4.` Use **df2** in the below cells to answer the quiz questions related to **Quiz 4** in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

In [14]:
```
prob_convert = df2.query('converted == 1').shape[0] / df2.shape[0]
prob_convert
```

Out[14]:
```
0.11959708724499628
```

b. Given that an individual was in the `control` group, what is the probability they converted?

In [15]:
```
control_population = df2.query('group == "control"')
control_converted = df2.query('converted == 1 and group == "control"')

(control_converted.shape[0] / control_population.shape[0])
```

Out[15]:
```
0.1203863045004612
```

c. Given that an individual was in the `treatment` group, what is the probability they converted?

In [16]:
```
treatment_population = df2.query('group == "treatment"')
treatmetn_converted = control_converted = df2.query('converted == 1 and group ==
"treatment"')

(treatmetn_converted.shape[0] / treatment_population.shape[0])
```

Out[16]:
```
0.11880806551510564
```

d. What is the probability that an individual received the new page?

In [17]:
```
#For further analysis, we are querying now individual who received old page
received_new_page = df2.query('landing_page == "new_page"')
received_old_page = df2.query('landing_page == "old_page"')


received_new_page.shape[0] / df2.shape[0]
```

Out[17]:
```
0.5000619442226688
```

e. Consider your results from a. through d. above, and explain below whether you think there is sufficient evidence to say that the new treatment page leads to more conversions.

Once control group converting rate is higher than treatment group, we fail to reject the null hypothesis.

Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

`1.` For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of p_{old} and p_{new}, which are the converted rates for the old and new pages.

$$H_0 : p_{old} >= p_{new}$$

$$H_1 : p_{old} < p_{new}$$

`2.` Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

Use a sample size for each page equal to the ones in **ab_data.csv**.

Perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use **Quiz 5** in the classroom to make sure you are on the right track.

a. What is the **convert rate** for p_{new} under the null?

In [18]:
```
p_new = df2['converted'].mean()
p_new
```

Out[18]:
```
0.11959708724499628
```

b. What is the **convert rate** for p_{old} under the null?

In [19]:
```
p_old = df2['converted'].mean()
p_old
```

Out[19]:
```
0.11959708724499628
```

c. What is n_{new}?

In [20]:
```
n_new = received_new_page.shape[0]
n_new
```

Out[20]:
```
145310
```

d. What is n_{old}?

In [21]:
```
n_old = received_old_page.shape[0]
n_old
```

Out[21]:
```
145274
```

e. Simulate n_{new} transactions with a convert rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

In [22]:
```
np.random.binomial(n_new, p_new)
```

Out[22]:
```
17218
```

f. Simulate n_{old} transactions with a convert rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

In [23]:
```
np.random.binomial(n_old, p_old)
```

Out[23]:
```
17253
```

g. Find p_{new} - p_{old} for your simulated values from part (e) and (f).

In [24]:
```
p_new - p_old
```

Out[24]:
```
0.0
```

h. Simulate 10,000 p_{new} - p_{old} values using this same process similarly to the one you calculated in parts **a. through g.** above. Store all 10,000 values in a numpy array called **p_diffs**.

In [25]:
```
p_diffs = []

for _ in range(int(1e5)):
    new_page_converted = np.random.binomial(n_new,p_new)
    old_page_converted = np.random.binomial(n_old, p_old)
    p_diff = new_page_converted/n_new - old_page_converted/n_old
    p_diffs.append(p_diff)
```

i. Plot a histogram of the **p_diffs**. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

In [26]:
```
plt.title('Simulated p-values')
plt.ylabel('Frequency')
plt.xlabel('p-values')
plt.hist(p_diffs);
```

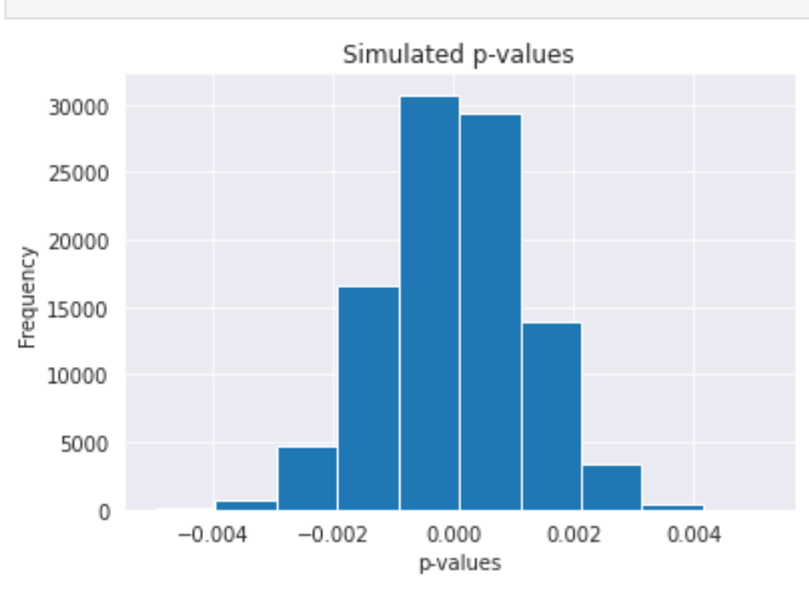

j. What proportion of the **p_diffs** are greater than the actual difference observed in **ab_data.csv**?

In [27]:
```
p_diff_orig = (treatment_population['converted'].mean() -
control_population['converted']).mean()
p_diff_orig
```

Out[27]:
```
-0.0015782389853555795
```

In [28]:
```
p_diffs = np.array(p_diffs)
p_diff_proportion = (p_diff_orig < p_diffs).mean()
print('proportion of p_diffs greater than p_diffs from ab_data.csv :
',p_diff_proportion)
```

```
proportion of p_diffs greater than p_diffs from ab_data.csv :  0.90443
```

k. In words, explain what you just computed in part **j.** What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

A p-value is the probability of observing your statistic if the null hypothesis is true.

The null hypothesis was that the difference in means would be equal or less than 0, and the alternative was the difference would be greater than 0. However, the difference is less than zero, and the p-value is very large. We do not have evidence to rejust the null.

l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer the the number of rows associated with the old page and new pages, respectively.

In [29]:
```
import statsmodels.api as sm

convert_old = df2.query('landing_page == "old_page" & converted == True').shape[0]
convert_new = df2.query('landing_page == "new_page" & converted == True').shape[0]

#These value was already calculated above
n_old = n_old
n_new = n_new
```

m. Now use `stats.proportions_ztest` to compute your test statistic and p-value. Here is a helpful link on using the built in.

In [30]:
```
stats, p_value = sm.stats.proportions_ztest([convert_old, convert_new], [n_old,
n_new], alternative='smaller')
print('stats: ', stats)
print('p_value: ', p_value)
```

```
stats:  1.3109241984234394
p_value:  0.9050583127590245
```

n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts **j.** and **k.**?

In [31]:
```
from scipy.stats import norm

#calculate the critical z_term
z_critical = norm.ppf(1-(0.05))
print('Zscore: ', stats)
print('Critical Zscore: ', z_critical)
```

```
Zscore:  1.3109241984234394
Critical Zscore:  1.6448536269514722
```

Since both p-values are equal and larger than 5% (0.05), and our Z-score are smaller than critical z-score, then we fail to reject the null, that is: old page is more efficient to convert than new page.

Part III - A regression approach

`1.` In this final part, you will see that the result you acheived in the previous A/B test can also be acheived by performing regression.

a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

Since the value we are trying to predict is categorical (convert or non converted), we should use a logistic regression

b. The goal is to use **statsmodels** to fit the regression model you specified in part **a.** to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create a column for the intercept, and create a dummy variable column for which page each user received. Add an **intercept** column, as well as an **ab_page** column, which is 1 when an individual receives the **treatment** and 0 if **control**.

In [32]:
```
df2['intercept'] = 1
df2[['ab_page', 'control']] = pd.get_dummies(df['group'])
df2.drop('control', axis=1, inplace=True)
df2.head()
```

Out[32]:

	user_id	timestamp	group	landing_page	converted	intercept	ab_page
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	1	1
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	1	1
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	1	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	1	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	1	1

c. Use **statsmodels** to import your regression model. Instantiate the model, and fit the model using the two columns you created in part **b.** to predict whether or not an individual converts.

In [33]:
```
import statsmodels.api as sm

logit = sm.Logit(df['converted'], df[['intercept', 'ab_page']])
results = logit.fit()
```

```
Optimization terminated successfully.
         Current function value: 0.366118
         Iterations 6
```

d. Provide the summary of your model below, and use it as necessary to answer the following questions.

In [34]:
```
results.summary()
```

Out[34]:

Logit Regression Results

Dep. Variable:	converted	**No. Observations:**	290584
Model:	Logit	**Df Residuals:**	290582
Method:	MLE	**Df Model:**	1
Date:	Mon, 28 Dec 2020	**Pseudo R-squ.:**	8.077e-06
Time:	00:10:47	**Log-Likelihood:**	-1.0639e+05
converged:	True	**LL-Null:**	-1.0639e+05
Covariance Type:	nonrobust	**LLR p-value:**	0.1899

	coef	std err	z	P>|z|	[0.025	0.975]
intercept	-2.0038	0.008	-247.146	0.000	-2.020	-1.988
ab_page	0.0150	0.011	1.311	0.190	-0.007	0.037

e. What is the p-value associated with **ab_page**? Why does it differ from the value you found in **Part II**?

Hint: What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in the **Part II**?

Answer
The p-value now is 0.19. The values is different from part II (0.90) because the hypotheses were different too. In part II, the null hyphotesis was new page are better to convert than old pages.

f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

Answer.
Adding more terms will increase model complexity, because can affect others variables, but can increase R-squared. One of the feature that worth to try add is duration that individual spent on the page, peharps time of day can be usefull.

g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives. You will need to read in the **countries.csv** dataset and merge together your datasets on the approporiate rows. Here are the docs for joining tables.

Does it appear that country had an impact on conversion? Don't forget to create dummy variables for these country columns - **Hint: You will need two columns for the three dummy variables.** Provide the statistical output as well as a written response to answer this question.

In [35]:
```
countries_df = pd.read_csv('./countries.csv')
df_new = countries_df.set_index('user_id').join(df2.set_index('user_id'), how='inner')
```

In [36]:
```
df_new['country']
```

Out[36]:
```
user_id
834778    UK
928468    US
822059    UK
711597    UK
710616    UK
          ..
653118    US
878226    UK
799368    UK
655535    CA
934996    UK
Name: country, Length: 290584, dtype: object
```

In [37]:
```
df_new[['CA', 'UK', 'US']] = pd.get_dummies(df_new['country'])
df_new[['ab_page', 'control']] = pd.get_dummies(df_new['group'])

#Let's use US as baseline country, so we don't need that column, as well as control.
df_new.drop(['US', 'control'], axis=1, inplace=True)
df_new
```

Out[37]:

	country	timestamp	group	landing_page	converted	intercept	ab_page	CA	UK
user_id									
834778	UK	2017-01-14 23:08:43.304998	control	old_page	0	1	1	0	1
928468	US	2017-01-23 14:44:16.387854	treatment	new_page	0	1	0	0	0
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	1	1	0	0	1
711597	UK	2017-01-22 03:14:24.763511	control	old_page	0	1	1	0	1
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	0	1	0	0	1
...	...	...	...	...	...	...	...	...	...
653118	US	2017-01-09 03:12:31.034796	control	old_page	0	1	1	0	0
878226	UK	2017-01-05 15:02:50.334962	control	old_page	0	1	1	0	1
799368	UK	2017-01-09 18:07:34.253935	control	old_page	0	1	1	0	1
655535	CA	2017-01-09 13:30:47.524512	treatment	new_page	0	1	0	1	0
934996	UK	2017-01-09 00:30:08.377677	control	old_page	0	1	1	0	1

290584 rows × 9 columns

h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model.

Provide the summary results, and your conclusions based on the results.

In [38]:
```
#Creating additional column with country and page column
df_new['CA_page'] = df_new['CA'] * df_new['ab_page']
df_new['UK_page'] = df_new['UK'] * df_new['ab_page']

df_new.head()
```

Out[38]:

	country	timestamp	group	landing_page	converted	intercept	ab_page	CA	UK	CA_page
user_id										
834778	UK	2017-01-14 23:08:43.304998	control	old_page	0	1	1	0	1	
928468	US	2017-01-23 14:44:16.387854	treatment	new_page	0	1	0	0	0	
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	1	1	0	0	1	
711597	UK	2017-01-22 03:14:24.763511	control	old_page	0	1	1	0	1	
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	0	1	0	0	1	

In [39]:
```
### Fit Your Linear Model And Obtain the Results
logit = sm.Logit(df_new['converted'], df_new[['intercept', 'CA', 'CA_page', 'UK',
'UK_page', 'ab_page']])
results = logit.fit()
results.summary()
```

```
Optimization terminated successfully.
         Current function value: 0.366109
         Iterations 6
```

Out[39]:

Logit Regression Results

Dep. Variable:	converted	**No. Observations:**	290584
Model:	Logit	**Df Residuals:**	290578
Method:	MLE	**Df Model:**	5
Date:	Mon, 28 Dec 2020	**Pseudo R-squ.:**	3.482e-05
Time:	00:10:51	**Log-Likelihood:**	-1.0639e+05
converged:	True	**LL-Null:**	-1.0639e+05
Covariance Type:	nonrobust	**LLR p-value:**	0.1920

	coef	std err	z	P>|z|	[0.025	0.975]
intercept	-2.0070	0.010	-207.045	0.000	-2.026	-1.988
CA	-0.0644	0.038	-1.679	0.093	-0.140	0.011
CA_page	0.0469	0.054	0.872	0.383	-0.059	0.152
UK	0.0257	0.019	1.363	0.173	-0.011	0.063
UK_page	-0.0314	0.027	-1.181	0.238	-0.084	0.021
ab_page	0.0206	0.014	1.505	0.132	-0.006	0.047

We still fail to reject null hyphotesis since convertion rate is similar. Also these interactions between countries and pages has larger p-values and could be decreasing significance of countries variables. As reported before, add duration of user interaction would be intesting, but since in this dataset we just have timestamp, which means only the moment that user clicked on the page, we can't get that information.

Let's take a look in duration of test in dataset to see if we can discover some potencial inside:

In [40]:
```
#Converting timestamp from object to datetime format
df_new['timestamp'] = pd.to_datetime(df_new['timestamp'])
df_new['timestamp'] = df_new['timestamp'].dt.date
print('Duration of test:', df_new['timestamp'].max() - df_new['timestamp'].min())
```

```
Duration of test: 22 days, 0:00:00
```

In [67]:
```
df_new['timestamp'].value_counts().sort_index()
```

Out[67]:
```
2017-01-02     5712
2017-01-03    13208
2017-01-04    13119
2017-01-05    12932
2017-01-06    13353
2017-01-07    13213
2017-01-08    13387
2017-01-09    13242
2017-01-10    13350
2017-01-11    13361
2017-01-12    13159
2017-01-13    13060
2017-01-14    13148
2017-01-15    13263
2017-01-16    13136
2017-01-17    13155
2017-01-18    13085
2017-01-19    13130
2017-01-20    13213
2017-01-21    13309
2017-01-22    13265
2017-01-23    13349
2017-01-24     7435
Name: timestamp, dtype: int64
```

We have just 22 days of data, with traffic well distributed. As sugestion we can collect 7 more days of data, to complete a month and see if the difference still the same.

Conclusions

Congratulations on completing the project!

Gather Submission Materials

Once you are satisfied with the status of your Notebook, you should save it in a format that will make it easy for others to read. You can use the **File -> Download as -> HTML (.html)** menu to save your notebook as an .html file. If you are working locally and get an error about "No module name", then open a terminal and try installing the missing module using `pip install <module_name>` (don't include the "<" or ">" or any words following a period in the module name).

You will submit both your original Notebook and an HTML or PDF copy of the Notebook for review. There is no need for you to include any data files with your submission. If you made reference to other websites, books, and other resources to help you in solving tasks in the project, make sure that you document them. It is recommended that you either add a "Resources" section in a Markdown cell at the end of the Notebook report, or you can include a `readme.txt` file documenting your sources.

Submit the Project

When you're ready, click on the "Submit Project" button to go to the project submission page. You can submit your files as a .zip archive or you can link to a GitHub repository containing your project files. If you go with GitHub, note that your submission will be a snapshot of the linked repository at time of submission. It is recommended that you keep each project in a separate repository to avoid any potential confusion: if a reviewer gets multiple folders representing multiple projects, there might be confusion regarding what project is to be evaluated.

It can take us up to a week to grade the project, but in most cases it is much faster. You will get an email once your submission has been reviewed. If you are having any problems submitting your project or wish to check on the status of your submission, please email us at dataanalyst-project@udacity.com. In the meantime, you should feel free to continue on with your learning journey by beginning the next module in the program.

In [ ]:
```

```